

Language Technology

<http://cs.lth.se/edan20/>
Self-Attention and the Transformer

Pierre Nugues

Pierre.Nugues@cs.lth.se
http://cs.lth.se/pierre_nugues/

September 28 and October 5, 2023



Transformers

After feedforward and recurrent networks, transformers are a third form of networks (in fact a kind of feedforward):

- An architecture proposed in 2017 based on the concept of **attention**
- Consists of a smart pipeline of matrices
- Transformers shown in this lecture, encoders, are trained with a **masked language model**
- They can learn complex lexical relations



Using Transformers

Goals of transformers:

- Encapsulate a massive amount of knowledge.
- In consequence trained on very large corpora
- Sometimes marketed as the ImageNet moment (See <https://runder.io/nlp-imagenet/>)

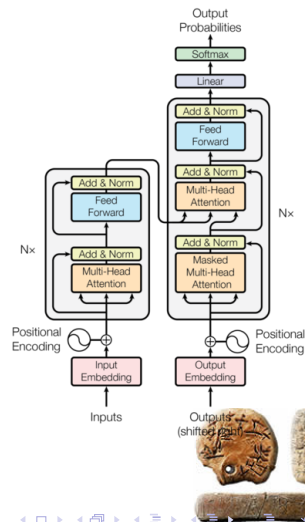
Transformers in practice:

- Large companies train transformers on colossal corpora, the pretrained models, requiring huge computing resources (<https://arxiv.org/pdf/1906.02243.pdf>)
- Mere users:
 - Reuse the models in applications
 - Fine-tune some parameters in the downstream layers



The Concept of Attention

- Reference paper: *Attention Is All You Need* by Vaswani et al (2017)
Link: <https://arxiv.org/pdf/1706.03762.pdf>
- Architecture consisting of two parts: encoder (left) and decoder (right)
- Implementation in PyTorch:
<https://nlp.seas.harvard.edu/2018/04/03/attention.html>
- We first consider the encoding part



Contextual Embeddings

Embeddings we have seen so far do not take the context into account
Attention is a way to make them aware of the context.

Consider the sentence:

I must go back to my ship and to my crew
Odyssey, book I

The word *ship* can be a verb or a noun with different meanings, but has only one GloVe embedding vector

Self-attention will enable us to compute contextual word embeddings.



Self-Attention


In the paper *Attention is all you need*, Vaswani et al. (2017) use three kinds of vectors, queries, keys, and values. Here we will use one type corresponding to GloVe embeddings.


We compute the attention this way:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V,$$

where d_k is the dimension of the input.

The softmax function is defined as:

$$\text{softmax}(x_1, x_2, \dots, x_j, \dots, x_n) = \left(\frac{e^{x_1}}{\sum_{i=1}^n e^{x_i}}, \frac{e^{x_2}}{\sum_{i=1}^n e^{x_i}}, \dots, \frac{e^{x_j}}{\sum_{i=1}^n e^{x_i}}, \dots, \frac{e^{x_n}}{\sum_{i=1}^n e^{x_i}} \right)$$




The meaning of QK^T

QK^T is the dot product of the GloVe vectors. It will tell us the similarity between the words

This is analogous to cosine similarity:

| | i | must | go | back | to | my | ship | and | to | my | crew |
|------|------|------|------|------|------|------|------|------|------|------|------|
| i | 1.00 | 0.75 | 0.86 | 0.76 | 0.73 | 0.90 | 0.35 | 0.65 | 0.73 | 0.90 | 0.42 |
| must | 0.75 | 1.00 | 0.85 | 0.68 | 0.87 | 0.69 | 0.42 | 0.69 | 0.87 | 0.69 | 0.45 |
| go | 0.86 | 0.85 | 1.00 | 0.84 | 0.84 | 0.81 | 0.41 | 0.68 | 0.84 | 0.81 | 0.49 |
| back | 0.76 | 0.68 | 0.84 | 1.00 | 0.83 | 0.76 | 0.49 | 0.77 | 0.83 | 0.76 | 0.51 |
| to | 0.73 | 0.87 | 0.84 | 0.83 | 1.00 | 0.68 | 0.54 | 0.86 | 1.00 | 0.68 | 0.51 |
| my | 0.90 | 0.69 | 0.81 | 0.76 | 0.68 | 1.00 | 0.38 | 0.63 | 0.68 | 1.00 | 0.44 |
| ship | 0.35 | 0.42 | 0.41 | 0.49 | 0.54 | 0.38 | 1.00 | 0.46 | 0.54 | 0.38 | 0.78 |
| and | 0.65 | 0.69 | 0.68 | 0.77 | 0.86 | 0.63 | 0.46 | 1.00 | 0.86 | 0.63 | 0.49 |
| to | 0.73 | 0.87 | 0.84 | 0.83 | 1.00 | 0.68 | 0.54 | 0.86 | 1.00 | 0.68 | 0.51 |
| my | 0.90 | 0.69 | 0.81 | 0.76 | 0.68 | 1.00 | 0.38 | 0.63 | 0.68 | 1.00 | 0.44 |
| crew | 0.42 | 0.45 | 0.49 | 0.51 | 0.51 | 0.44 | 0.78 | 0.49 | 0.51 | 0.44 | 1.00 |



Vaswani's attention score

The attention scores are scaled and normalized by the softmax function.

$$\text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right),$$

| | i | must | go | back | to | my | ship | and | to | my | crew |
|------|------|------|------|------|------|------|------|------|------|------|------|
| i | 0.36 | 0.05 | 0.07 | 0.05 | 0.04 | 0.19 | 0.01 | 0.02 | 0.04 | 0.19 | 0.01 |
| must | 0.14 | 0.20 | 0.10 | 0.06 | 0.11 | 0.10 | 0.03 | 0.05 | 0.11 | 0.10 | 0.02 |
| go | 0.18 | 0.09 | 0.14 | 0.09 | 0.08 | 0.13 | 0.02 | 0.04 | 0.08 | 0.13 | 0.02 |
| back | 0.14 | 0.05 | 0.09 | 0.19 | 0.08 | 0.12 | 0.03 | 0.06 | 0.08 | 0.12 | 0.03 |
| to | 0.11 | 0.11 | 0.09 | 0.09 | 0.15 | 0.08 | 0.04 | 0.07 | 0.15 | 0.08 | 0.03 |
| my | 0.19 | 0.03 | 0.05 | 0.04 | 0.03 | 0.29 | 0.01 | 0.02 | 0.03 | 0.29 | 0.01 |
| ship | 0.03 | 0.03 | 0.03 | 0.04 | 0.05 | 0.03 | 0.55 | 0.03 | 0.05 | 0.03 | 0.13 |
| and | 0.10 | 0.08 | 0.07 | 0.10 | 0.12 | 0.09 | 0.04 | 0.15 | 0.12 | 0.09 | 0.04 |
| to | 0.11 | 0.11 | 0.09 | 0.09 | 0.15 | 0.08 | 0.04 | 0.07 | 0.15 | 0.08 | 0.03 |
| my | 0.19 | 0.03 | 0.05 | 0.04 | 0.03 | 0.29 | 0.01 | 0.02 | 0.03 | 0.29 | 0.01 |
| crew | 0.06 | 0.05 | 0.05 | 0.06 | 0.05 | 0.06 | 0.21 | 0.04 | 0.05 | 0.06 | 0.31 |



Attention

We use these scores to compute the attention.

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V,$$

For *ship*:

```
attention_ship = (0.03 * embeddings_dict['i'] +  
                  0.03 * embeddings_dict['must'] +  
                  0.03 * embeddings_dict['go'] +  
                  0.03 * embeddings_dict['back'] +  
                  0.04 * embeddings_dict['to'] +  
                  0.05 * embeddings_dict['my'] +  
                  0.55 * embeddings_dict['ship'] +  
                  0.03 * embeddings_dict['and'] +  
                  0.05 * embeddings_dict['to'] +  
                  0.03 * embeddings_dict['my'] +  
                  0.13 * embeddings_dict['crew'])
```

where the *ship* vector received 13% of its value from *crew*



Code Example

Experiment: Jupyter Notebook: `ch12_attention.ipynb`
(First part of the notebook)



Multihead Attention

This attention is preceded by dense layers:

If X represents complete input sequence (all the tokens), we have:


$$Q = XW_Q,$$

$$K = XW_K,$$

$$V = XW_V.$$

And followed by another dense layer.

In addition, most architectures have parallel attentions, where the outputs (called heads) are concatenated (multihead)



From *Attention is all you need*,
Vaswani et al. (2017)



Code Example

Keras has an implementation of this architecture with the `MultiHeadAttention()` layer.

Experiment: Jupyter Notebook: `ch12_attention.ipynb`
(Second part of the notebook)



Transformers: The Encoder

In transformers, the encoder is a structure, where:

- 1 The first part of the layer is a multihead attention;
- 2 We reinject the input to the attention output in the form of an addition:

$$X + \text{Attention}(Q, K, Q).$$


This operation is called a skip or residual connection, which improves stability.

- 3 The result is then normalized per instance, i.e. a unique sequence, defined as:

$$x_{i,j_{norm}} = \frac{x_{i,j} - \bar{x}_{i,.}}{\sigma_{x_{i,.}}}.$$

The input distribution is more stable and improves the convergence

- 4 It is followed by dense layers.



Left part, from *Attention is all you need*, Vaswani et al. (2017)



Code Example

Experiment: Jupyter Notebook: [https://github.com/fchollet/deep-learning-with-python-notebooks/](https://github.com/fchollet/deep-learning-with-python-notebooks/blob/master/chapter11_part03_transformer.ipynb)

[blob/master/chapter11_part03_transformer.ipynb](https://github.com/fchollet/deep-learning-with-python-notebooks/blob/master/chapter11_part03_transformer.ipynb) from Chollet's book, first part up to Using the Transformer encoder for text classification

As a personal work and to gain a deeper understanding, you can read a tutorial in PyTorch:

<https://nlp.seas.harvard.edu/2018/04/03/attention.html>



Training Transformers

Transformers, such as BERT, are often trained on masked language models with two tasks:

- 1 For a sentence, predict masked words: We replace 15% of the tokens with a specific mask token and we train the model to predict them. This is just a cloze test;
- 2 For a pair of sentences, predict if the second one is the successor of the first one;

Taking the two first sentences from the *Odyssey*:

*Tell me, O Muse, of that ingenious hero who travelled far and wide after he had sacked the famous town of Troy.
Many cities did he visit, and many were the nations with whose manners and customs he was acquainted;*



Masked language models

We add two special tokens: [CLS] at the start of the first sentence and [SEP] at the end of both sentences, and the token [MASK] to denote the words to predict.

We would have for the first task:

[CLS] Tell me, O Muse, of that [MASK] hero who travelled far and wide [MASK] he had sacked the [MASK] town of Troy. [SEP]

For the second task, we would have as input:

[CLS] Tell me, O Muse, of that [MASK] hero who travelled far and wide [MASK] he had sacked the [MASK] town of Troy. [SEP] Many cities did he [MASK visit], and many were the [MASK nations] with whose manners [MASK and] customs he was acquainted; [SEP]

where the prediction would return that the second sentence is the next one (as opposed to random sequences)



Positional Embeddings

BERT (the first transformer, Devlin et al. (2019)) maps each token to three embedding vectors:

- the token embedding,
- the position of the token in the sentence (positional embeddings), and
- the segment embeddings (we will skip this part).

The three kinds of embeddings are learnable vectors.

In the BERT base version, each embedding vector has 768 dimensions.

Let us consider two sentences simplified from the *Odyssey*:

Tell me of that hero. Many cities did he visit.

| Input: | [CLS] | Tell | me | of | that | hero | [SEP] | Many | cities | did | he | visit | [SEP] |
|----------|-------------|-------------------|-----------------|-----------------|-------------------|-------------------|-------------|-------------------|---------------------|------------------|-----------------|--------------------|-------------|
| Token | $E_{[CLS]}$ | E_{tell} | E_{me} | E_{of} | E_{that} | E_{hero} | $E_{[SEP]}$ | E_{many} | E_{cities} | E_{did} | E_{he} | E_{visit} | $E_{[SEP]}$ |
| Segment | E_A | E_A | E_A | E_A | E_A | E_A | E_A | E_B | E_B | E_B | E_B | E_B | E_B |
| Position | E_0 | E_1 | E_2 | E_3 | E_4 | E_5 | E_6 | E_7 | E_8 | E_9 | E_{10} | E_{11} | E_{12} |



Code Example

Experiment: Jupyter Notebook: [https:](https://github.com/fchollet/deep-learning-with-python-notebooks/blob/master/chapter11_part03_transformer.ipynb)

[//github.com/fchollet/deep-learning-with-python-notebooks/blob/master/chapter11_part03_transformer.ipynb](https://github.com/fchollet/deep-learning-with-python-notebooks/blob/master/chapter11_part03_transformer.ipynb) from Chollet's book, second part from Implementing positional embedding as a subclassed layer



Model size

Transformers are trained on large corpora like the colossal clean crawled corpus (<https://arxiv.org/abs/2104.08758>) and encapsulate semantics found in text in the form of numerical matrices.

This results in large models (Devlin et al., 2019):

In this work, we denote the number of layers (i.e., Transformer blocks) as L , the hidden size as H , and the number of self-attention heads as A . We primarily report results on two model sizes: $BERT_{BASE}$ ($L=12$, $H=768$, $A=12$, Total Parameters=110M) and $BERT_{LARGE}$ ($L=24$, $H=1024$, $A=16$, Total Parameters=340M).


Transformers can then act as pre-trained models for a variety of tasks. See the list from Huggingface

Finally, an interesting reading: <https://sayakpaul.medium.com/an-interview-with-colin-raffel-research-scientist-at-google-5>



Applying Transformers

- Matrices in the transformer architecture encapsulate massive knowledge from text.
- We can apply them to tasks beyond what they have been trained for (masked language model)
- We use them as pretrained models and fine-tune a so-called dedicated *head*.
- The simplest head is a logistic regression

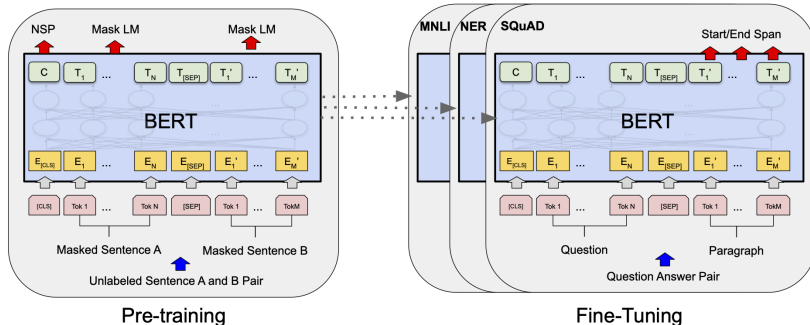


Picture from Wolf et al., Transformers: State-of-the-art Natural Language Processing, EMNLP Demos 2020.



Transfer Learning


Transfer learning consists of a costly **pretraining** step and an adaptation to applications called **fine-tuning**.



from Devlin et al., *BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding*, 2019. Note the picture comes from the second version of the paper from 2019



Application: Sentence Classification



(b) Single Sentence Classification Tasks:
SST-2, CoLA

from Devlin et al., *BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding*, 2019



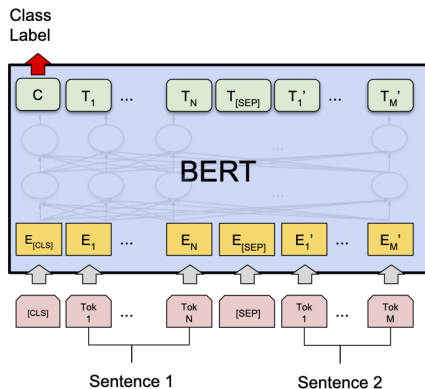
Code Example

Experiment: Jupyter Notebook: [https:](https://github.com/fchollet/deep-learning-with-python-notebooks/blob/master/chapter11_part03_transformer.ipynb)

[//github.com/fchollet/deep-learning-with-python-notebooks/blob/master/chapter11_part03_transformer.ipynb](https://github.com/fchollet/deep-learning-with-python-notebooks/blob/master/chapter11_part03_transformer.ipynb) from Chollet's book, first part from Using the Transformer encoder for text classification



Application: Sentence Pair Classification

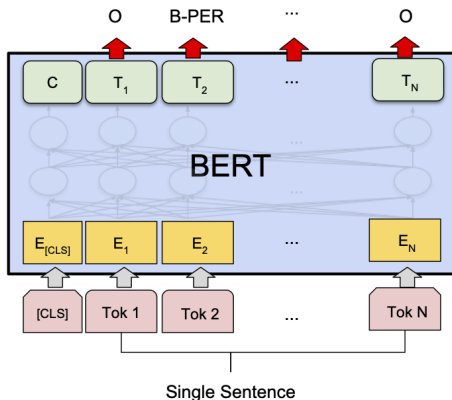


(a) Sentence Pair Classification Tasks:
MNLI, QQP, QNLI, STS-B, MRPC,
RTE, SWAG

from Devlin et al., *BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding*, 2019



Application: Sequence Tagging



(d) Single Sentence Tagging Tasks:
CoNLL-2003 NER

from Devlin et al., *BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding*, 2019



Stanford Question Answering Dataset (SQuAD)

- Consists of 100,000 questions and paragraphs from wikipedia containing the answers
- The answer is a segment in the text (factoid QA)
- Complemented by SQuAD 2.0 with unanswerable questions
- SQuAD started an intense competition. See the impressive leaderboard

<https://rajpurkar.github.io/SQuAD-explorer/>

Form Rajpurkar et al., *SQuAD: 100,000+ Questions for Machine Comprehension of Text*, 2016

In meteorology, precipitation is any product of the condensation of atmospheric water vapor that falls under **gravity**. The main forms of precipitation include drizzle, rain, sleet, snow, **grau-pel** and hail... Precipitation forms as smaller droplets coalesce via collision with other rain drops or ice crystals **within a cloud**. Short, intense periods of rain in scattered locations are called "showers".

What causes precipitation to fall?

gravity

What is another main form of precipitation besides drizzle, rain, snow, sleet and hail?

grau-pel


Where do water droplets collide with ice crystals to form precipitation?

within a cloud

Figure 1: Question-answer pairs for a sample passage in the SQuAD dataset. Each of the answers is a segment of text from the passage.



Application: Question Answering



(c) Question Answering Tasks:
SQuAD v1.1

from Devlin et al., *BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding*, 2019

